

Compute Servers for Teaching Big Data

Anonymous- for review

May 21, 2025

keywords: big data, cloud computing, data science, education, pedagogy, statistical computing

Abstract

We compare and contrast a variety of compute server options for teaching an undergraduate statistics class centered on analyzing large, real world data sets. These server options include institutionally supported servers, personal physical servers, public cloud services, and commercial cloud services. We briefly describe the course, share our experience with each of these options, advantages and disadvantages, some technical considerations, and suggest when to choose which option. Looking to the future, we outline requirements for a broadly shared educational compute service that will allow more instructors to teach similar classes without the overhead of administering their own server.

1 Introduction

Imagine yourself as a student who has just graduated college with an undergraduate degree in applied math or statistics, and started their first job working with data. The new manager gives them a challenge: “Here are your login credentials for our cloud provider services. Here are a few hundred gigabytes of messy data that nobody has yet had time to look at. Familiarize yourself with it, and find us some insights.” This is certainly a tall order, but the graduate who has the ability to even *begin* such a task is valuable and rare indeed. Success requires technical computing skills, sound statistical reasoning, and an ability to ‘think with data’.

The implicit goal of STAT 129, a class in ‘Big Data’ at California State University, Sacramento, is to prepare students for the scenario described above. This paper contains some of the key insights and experience from developing and teaching this

course for the past five years. Section 2 describes the class, including what kind of data suits our pedagogical purpose and why we need a compute server. Section 3 provides an overview of the role of the server in our class, and discusses the two physical server options: institutionally supported and personal. Section 4 covers cloud server options, including publicly funded and commercial. Section 5 concludes with our ‘wish list’ for an ideal shared educational compute service that will allow many instructors to simultaneously teach similar classes without the overhead of administering their own server.

2 The class: STAT 129

STAT 129, titled ‘Analyzing and Processing Big Data’, is an upper division elective course offered by the Mathematics and Statistics Department at California State University, Sacramento, and is required for students pursuing a bachelor’s degree in statistics with an emphasis in data science. The catalog course description is:

Statistical analysis of large, complex data sets. Topics include memory efficient data processing, the split-apply-combine strategy, rewriting programs for scalability, handling complex data formats, and applications such as statistical learning, dimension reduction, and efficient data representation. Students will access data and run code on remote servers.

STAT 129 uses technologies including bash, Python, and SQL to analyze large data sets, with the intention of exposing students to common tools in data science. Students tend to be highly motivated, working towards a job in data science after

graduation. STAT 129 requires a prerequisite class in programming, and the most successful students came into the class with strong programming skills.

The philosophy behind STAT 129 is to provide students with real world data sets and ask open ended questions, following in the footsteps of Deborah Nolan and Duncan Temple Lang [6]. We are statistics faculty with a background in computing, and we’ve developed and taught STAT 129 each Spring semester since 2021. Enrollment is capped at 25 students, and STAT 129 has enjoyed strong enrollment, with 18-25 students per year for the last three years. Feedback from our data forward approach has been overwhelmingly positive, with several alumni reporting STAT 129 as the most practical class they have taken in college.

2.1 Which Data?

Conventionally, big data refers to large, unwieldy data sets that require some effort to analyze [7]. For the purposes of STAT 129, ‘big data’ has a precise meaning: a data set that is large enough such that one cannot load the entire data set into a laptop’s memory at once. In a typical introductory data analysis class, a student loads a single data file into memory and proceeds with the analysis steps [12]. This process works great if the data fits into memory, but if the data does not fit into memory, then it does not work at all, because it fails at the loading step. Data sets that are larger than memory require entirely different approaches, such as first streaming through to filter out a subset of interest. Our goal is to choose data sets that are comfortably over this memory threshold, such that the techniques we are learning are motivated, but not gratuitously large such that answering a basic question takes

an excessive amount of time. Practically speaking, at the time of this writing, a typical student laptop will have between 4 and 32 GB of memory, and data sets that are on the order of 10-100 GB on disk work well to motivate the techniques. We assign data analysis tasks that typically take 1 to 10 minutes if done sequentially, and are an order of magnitude faster if done in parallel.

When teaching STAT 129, we look for data sets that satisfy the following criteria.

1. **Size:** The files should be large enough to motivate the computational techniques taught in class, but not excessively large, as described above.
2. **Publicly Available:** The data should be readily available for legal public download. Special permissions, non disclosure agreements, and licensing all take time away from teaching.
3. **Format:** The data storage format should be a common one that students will encounter in data science work, so that students can directly apply what they learn in class. Each semester of STAT 129, we work with at least three data sets: a plain text tabular format such as CSV, a nested format such as XML or JSON, and a SQL database. We spend several weeks working with unstructured text, which we usually extract as a field inside one of the structured formats above. The data is usually compressed, with one data set stored across multiple files. Avoid proprietary, novel, or highly specialized data formats.
4. **Background Knowledge:** A college student should be able to understand what the data actually represents in a single class period. For example, in the Global Historical Climatology Network-Daily Database [5], each row rep-

resents one weather observation for a specific date and location, which is easy to explain.

The above criteria do not necessarily apply to more domain specific classes. For example, a bioinformatics class studying genomes would do well to focus on the specialized formats of genomic data. Amazon Open Data [9] has provided a reliably good source of large data sets fitting these criteria. The US government [1] is another good source.

It may be possible to pre-process the data set before providing it to students to make it better suit the goals of the course. We often find interesting data sets that are close, but not quite what we are looking for. Consider the public tax return data on nonprofits provided by the Internal Revenue Service of the United States[8]. We removed the XML namespace information in these files so that students learning XPATH can write:

```
//ActivityOrMissionDesc
```

to find the ‘activity or mission description’ for a nonprofit rather than the cumbersome:

```
//*[local-name()='ActivityOrMissionDesc']
```

The latter is distracting for students just learning XPATH.

Large data sets are the heart of STAT 129. To properly analyze them we need to provide our students with a compute server.

2.2 Why a Server?

Students must access data and run code on remote servers for a class on ‘Big Data’, because that is where big data is typically found: on remote servers. There is a qualitative difference between running code on a laptop with a monitor that is sitting in front of you, and on a machine that you never actually see, that you can only access through `ssh`, secure shell, a purely text based way of interacting with a remote computer. After opening a terminal and logging in to the remote server with `ssh`, all you see is a `bash` prompt waiting for you to type commands; this is the command line. For example, the following shows what it would look like after the user Jane Doe logged in to the remote server via `ssh` and entered the `pwd` command to print her working directory:

```
jane@nsm-stats:~$ pwd
/home/jane
jane@nsm-stats:~$
```

Our goal with embracing the command line is to move students beyond the initial unfamiliarity and discomfort that users experience when they first move away from a Graphical User Interface (GUI). The command line is the only way for STAT 129 students to access the large data sets and complete the assignments, quizzes, and exams. Multiple alumni of STAT 129 have reported back that it was specifically the command line skills that led directly to them getting a variety of jobs.

A single server hits the sweet spot for STAT 129; clusters are less ideal, because of the overhead and the abstraction they introduce. One topic that we cover in STAT

129 is reading the output of `top/htop` to identify a bottleneck in a `bash` pipeline that is streaming hundreds of GB of data. It is possible to do this with a cluster, but it takes either more advanced knowledge, or tools specific to that system. Whenever possible we try to teach the tools that are the most common and widely applicable, to better serve the students in their future careers.

2.3 Student Experience

3 Servers

Which server option is best depends on the goals of the course, the level of institutional support, the number of students enrolled and their background, as well as the instructor's preferences. The overarching goal of STAT 129 is to statistically analyze real data sets, using professional software. Students use the server for every assignment and assessment for 15 weeks, so it makes sense to invest some time teaching them skills and tools with steeper learning curves, such as terminal based text editors. If a different class only used a server for two weeks of the semester to examine a single data set, then it probably makes more sense to use a more approachable web browser based graphical user interface to the server such as Jupyter Notebooks [4], which will inform the best choice of server options.

Here are some details of our specific situation. We had no servers available to use for teaching at the department, college, or university level; however, the college did generously provide us with startup funds, which funded the experimentation that led to this paper. The Mathematics and Statistics Department at Sacramento State

University plans to offer STAT 129 once a year for the foreseeable future, so we need a long term solution for a server. In a given year we will have one instructor, 15-30 students, and no staff, teaching assistants, or other support. We do not need a large machine for any purpose other than teaching. We have some professional background in programming, and our server administration experience has proven especially useful.

What follows are the server options, ordered by our personal preference. We have yet to discover a perfect solution, and what is best for you may well differ from our preferences.

3.1 Institutional Support

Institutional support for teaching a class like STAT 129 varies widely. If you are in the lucky position of having compute infrastructure with dedicated support in your department, college, or university, then this is an ideal option for teaching classes. You simply ask the system administrator for what you need, and they make it happen. The challenge for you as the instructor is determining what you need!

We have used a cluster with SLURM [13] job queueing software, common in high performance computing (HPC), for teaching Stat 141C (Big Data and High Performance Computing) at University of California, Davis in 2019. The system administrator helped us with the following tasks related to the class:

- Selecting the right servers to use for teaching.
- Creating user accounts from a student roster.

- Installing R, Python, along with various packages and libraries of software.
- Troubleshooting connectivity and user account issues. For teaching undergraduate courses, timely, helpful support is much more important than physical compute power.
- Setting up and configuring an open source Postgres database for students.

Of these tasks, setting up the Postgres database took the most time and expertise, and we were grateful to have the support. A system administrator can do many other things as well. If you do not have institutional support, then it is possible to be your own systems administrator, and we consider this situation next.

3.2 Personal Physical Server

A personal physical server is essentially the same as a powerful desktop machine, except that you can access it remotely, and it does not need a monitor. The machine itself might be in a rack in a server room on campus, or in someone's office. The server rack option is generally preferable, as the machines have more central processing units (CPU) and memory, better internet connectivity, and they do not clutter up your space. Figure 1 shows the physical machine in our office that is our current solution, after having explored and experienced the others described in this article. We find that it is the most pragmatic long term solution, requiring the least effort on our part.

The major advantages of the personal physical server are cost and simplicity. Regarding cost, we have found it easier to fund a one time purchase of a few thousand



Figure 1: A simple server under a desk, our current solution, works fine for teaching.

dollars, compared to finding a sponsor who will commit to the ongoing variable subscription fees of the cloud services, even if the total cost would be lower for the cloud subscription. Physical servers are simple because you can set it up one time, and then use the same machine for several years. We did not need any specialized knowledge beyond the basics. We completely own and administer the machine, so we can install whatever we want, and whenever we want. It is relatively easy and direct to automate certain grading tasks, for example, an instructor can use administrative privileges to look at `history` to see every command a student typed during a quiz. If the machine is in your office, then when you launch a big job, the

machine will heat up and the cooling fan will turn on, which provides a very real connection to the amount of data being processed for the students in office hours, which is fun.

The primary disadvantage of a single physical machine is that if it fails, then the class effectively stops. If students cannot use the server, then they cannot do their homework, and they cannot do quizzes. If you are managing the server, then you must identify and fix the problem. Students have made mistakes and filled up all the disk space, and we needed to diagnose and fix the issue on our own. Pedagogically, this mistake followed by the unusable server highlights the importance of not creating unnecessary copies when working with large amounts of data, and it makes students more aware of exactly what happens if they inadvertently uncompress and save an excessive amount of data. The fix is straightforward: just delete the big files. A more proactive solution to this common problem is to impose a quota on disk usage by the home directory. We are taking a risk that the server could fail in a more serious way, but it is a calculated risk.

The server does not need to be powerful or expensive. For a class of 15 students, we found that a single 32 core CPU, a GPU, and 128 GB of memory was just sufficient for students to experience the class assignments in the way we intended. This machine cost around \$3,000 when purchased in 2022. We always recommend that students do their homework well before the due date, because if a handful of users are simultaneously trying to use parallel jobs with all available CPU cores, then they will not see the speedup they expect. In practice, it happened occasionally that students who waited until the last hour to do an assignment did not get to see their

final result, but they were not surprised, because we had warned them many times and they knew to expect this. We have used a 16 core CPU machine in previous semesters, and found that it was not sufficient for 15 students. With more students, we believe that students would feel crowded on the server.

For a server in your office, you must provide remote access so that others can log on to the machine. Campus IT handled this by providing a DNS host name. Our students need to be either on campus or at home on the campus virtual private network (VPN), and they can login via `ssh`. We recommend against setting up a physical server at your home, because providing reliable remote access could be tough depending on your internet provider, and there is no reason you should pay for the electricity.

The physical server feels out of date, and lacks the cachet of being ‘the cloud’. Despite that, it provides a legitimate and simple way to get students using a remote server. From a student’s perspective, there is no difference between logging in via `ssh` to a physical machine, or to a virtual machine running on the cloud.

4 Cloud

The cloud offers the most flexibility of all the server options. Indeed, all the different possibilities can be overwhelming. In STAT 129, students process data by logging into a virtual machine (VM), a logical rather than a physical computer. We have experimented with two ways of providing students with cloud access: local user accounts for each student on a VM, and administrative cloud accounts for every

student. Both options have a place, depending on the goals of the class.

Administrative cloud accounts provide students with a richer cloud experience, as they create their own virtual machines which they have complete control over. They are the administrators to these machines with `sudo` access, so they can do anything they wish, and any mistakes they make will not affect anyone else. Students can create machines that are as large as they like, and delete them when they are done, which solves the problem of many students trying to simultaneously use limited resources on a single machine. Cloud accounts would be a good option if the course is designed to teach them more about infrastructure, for example, how to implement a full stack web application. This rich experience does come with some additional complexity and responsibility, for example, students need to load and save their work as they create and delete machines.

Local user accounts on a single VM allow students to focus more on the data analysis, rather than the cloud ideas and infrastructure. In this scenario, the instructor creates a VM that is large enough to accommodate all of the students simultaneously, and students access through `ssh` or a web interface. The downside here is that you need to pay for a sufficiently large machine all the time, or else manually load balance by switching to smaller or larger machines as the demand changes.

You can combine the simpler single VM user accounts together with the richer administrative cloud accounts. For example, we have started the semester with a single VM large enough for the class, which students all log in to and use for simple assignments and in class activities, and then students progress to setting up their own VM's later in the semester.

Now we turn our attention to the two cloud provider options: public and commercial.

4.1 Public Cloud

We have had an excellent experience teaching STAT 129 using the NSF funded cloud provider, Jetstream2 [3]. In this model, the instructor applies for an allocation, the reviewers award a certain number of compute hours, and the instructor sets up the accounts on the machines. We applied, were approved, and had everything set up in less than a week. The reviewer awarded 100,000 service units to teach the class, which was several times more than we requested, and this proved to be plenty for 20 students over a semester.

The Jetstream2 website [11] states:

Jetstream2 makes cutting-edge high-performance computing and software easy to use for research regardless of your project's scale, even if you have limited experience with supercomputing systems.

Cloud-based and on-demand, the Jetstream2 environment also includes discipline-specific apps. You can even create virtual machines that look and feel like your lab workstation or home machine—but with thousands of times the computing power.

The benefits of Jetstream2 are the free cost for students, the simplicity compared to commercial cloud providers, and the ability to connect seamlessly to large data sets. Thanks to NSF funding, our students, many of whom are low income, received

accounts without providing a credit card. The user interface shown in Figure 2 is simple and easy to use, with a manageable number of options.

Jetstream2 caters directly to the academic wishing to analyze large data sets through Manila Filesystems as a service. We followed the Jetstream2 documentation and mounted a 1 TB Manila file share to the VMs, which allowed students to simultaneously access and analyze shared data as if it was local to their machine [2]. In our experience, Jetstream2 has lived up to its claims of making cloud computing easy and accessible.




Figure 2: This screenshot shows how to create a new instance of a virtual machine using Jetstream2. This user interface is simple and easy to teach with.

The main disadvantages of Jetstream2 for teaching college classes are the security model and the lack of safeguards. Jetstream2 prioritizes the needs of individual researchers and small research groups, and they support this primary use case very well. Every user gets complete access to every part of the account on the allocation, which keeps things simple. However, if you give students user accounts so that they can have the experience of creating their own VM's, then they will have administra-

tive access to everything the instructor creates, and to every other student's machine. There are no limits on how much computational time a user can use, so it is possible for one user to use up all the credits for the allocation. In other words, the only thing that is stopping them from doing something bad is their own awareness and respect for others, which is not ideal for a group of 20-30 novices. We mostly got around these issues by setting very clear expectations for behavior, with the understanding that this whole class is based on trust. Even so, a student unintentionally left a large machine running for several weeks. We were not monitoring as diligently as we should have, and this used up a significant portion of our available compute time, but we still had plenty to get through the course.

Finally, our allocation on Jetstream2 required annual renewal, which makes us hesitate to rely on it as a long term solution for teaching STAT 129. We expect the reviewers would renew every year if there is funding available, but there are no guarantees. Annual renewal is not an issue if someone is paying for commercial cloud services, and we turn our attention here next.

4.2 Commercial Cloud

Several companies offer commercial cloud services. We focused mainly on Amazon Web Services (AWS), the biggest and most popular vendor, to provide our students with marketable experience for their resumes. Our favorite feature of AWS is the direct access to Amazon's high quality open data sets through the `aws s3` command line tools, which eliminate the step of downloading the data and making it available to students.

Sustainable ways to provide servers for teaching a class like STAT 129 are essential, as discussed in Section 3. A sustainable funding model is to make students pay for commercial cloud services, and this does not need to be expensive. In Spring 2022, we found that students on AWS can get a rich data processing experience for \$20 or less in one semester, which is less than a low cost textbook for a single course. These are commercial services, so a student must have a credit card and be prepared to pay for any compute time and storage that they use. If students make a mistake, say using a powerful VM and forgetting to shut it off, then they can easily run up an expensive and unexpected bill that can be tough for a college student to cover. In the past, AWS offered an educational program with free credits for students, which did not even require a credit card. We used this program, but they discontinued it, which left us looking for other free options.

Complexity and vendor specificity are problems when teaching with cloud services. Amazon’s EC2 (cloud compute server) allows students to create their own VM’s, and it “provides a wide selection of over 750 instance types optimized to fit different use cases.” [10] There are also hundreds of different services and many world regions to choose from. This is an overwhelming amount of choice for a student or even an instructor, and the available options change regularly. Furthermore, the services often have company specific names, which makes it difficult for an instructor to understand what they are. Indeed, the main purpose of the ‘AWS Cloud Practitioner’ certification seems to be to understand all of these proprietary names and services. As faculty at a public institution, our goal is to teach the overarching concepts and portable technologies that a student can take and apply anywhere, regardless of the

specific vendor.

Critiques aside, commercial cloud services still offer an appealing option for teaching STAT 129. With the commercial cloud you can get started very quickly with minimal work and investment. If you already have experience with a particular cloud provider, or you wish to build that experience from the ground up, then this is an opportunity. If your institution has some connection with a cloud provider, then it may be mutually beneficial to leverage those connections. For example, the company could provide cloud services for free in exchange for the access to recruit graduates into that company.

If the goal of your course is to focus more on a specific technology such as R, Python, or SQL, then you should consider popular vendors such as Posit Cloud (formerly RStudio), Anaconda Cloud, or Snowflake. These vendors use commercial cloud services as a foundation, abstract away many of the complicated details of server administration, and provide a familiar web based user interface. None of the available cloud services that we have explored or mentioned have proven to be ideal for teaching STAT 129. We conclude by describing what *would* be ideal, based on our experience.

5 Moving Forward

There are many technical details and considerations when relying on a server to teach an undergraduate data analysis class like STAT 129, and we have only provided a broad overview here. For a lone instructor with a heavy teaching load and no support,

the barriers to entry are significant. If it was easier to procure, configure, and run a server, then we believe that more instructors would be interested in teaching these courses.

We would welcome a shared solution, cloud based or otherwise, that provides this capability publicly to many instructors simultaneously teaching similar courses at other institutions. It should have the following features (listed alphabetically):

1. **data science tools** Students should be able to use common data science programming technologies such as R, Python, and some form of SQL. Web based tools such as Jupyter Notebooks certainly broaden the appeal and potential user base of the server, but they are not strictly necessary.
2. **free or low cost** Free for students is ideal. If free is not possible, then the cost over the semester should not exceed that of a reasonably priced textbook.
3. **instructor support** Instructors without a background in systems administration should be able to teach on the server, and receive some level of technical and community support. This is perhaps the hardest requirement, because the best way to provide dedicated support is through paid staff.
4. **long term sustainability** If instructors and institutions are going to invest in creating courses that rely on a server, then they need to have some assurance that the server will continue to exist. A 1 year pilot program might be a useful start, but it is much more compelling to have a guarantee that a server will be available for teaching for the next 5 years.

5. **minimal specificity** It is ideal if students in a one semester data science class do not need to learn concepts specific to one system or vendor, as discussed in Sections 2.2 and 4.2.
6. **novice user safeguards** It should not be possible for one student to delete or damage course infrastructure, either accidentally or on purpose. Standard Linux user accounts without administrative privileges work fine.
7. **parallel capability** The parallel processing capabilities of the server should exceed those of a common laptop to provide a compelling reason to actually use the server. In practice, with 10 CPU's students can experience an order of magnitude speedup compared to a serial program, and this is enough to motivate the techniques in STAT 129. An abundance of compute power is definitely not necessary if instructors follow the advice in Section 2.1.
8. **shared data storage** It should be easy for an entire class to simultaneously read and share the same large data sets. 100 GB of storage would suffice, 1 TB of storage would be great, and 10 TB would be more than enough.
9. **shell access** Students should be able to interact with their computer using bash so that they can use data processing pipelines, scripts, and version control.

The above instructor level requirements imply other requirements at the system level, and these require serious thought. For example, if a shared solution is going to simultaneously support 20 instructors teaching 400 students, then load balancing is a real concern. We leave the system level requirements to the professionals.

Fostering a culture of support, mentorship, and growth for instructors is more important than the logistics of setting up a server. We academics should value programming skills, and find ways to support instructors who may be teaching new and unfamiliar topics. Some specific things we can do are: encourage our peers to accept the challenge of teaching material that may be outside of their comfort zones, share our course materials, and meet with peers regularly when they are teaching this new material for the first time.

Large data sets are here to stay, and we serve our students well by teaching them how to actually work with large data. It is quite uncommon to see a course designed around the student experience on a server, particularly in a mathematics and statistics department. We hope that, in the future, compute servers become to a data science bachelor's degree program somewhat like physical laboratory space is to a chemistry program: a necessity.

References

- [1] United States General Services Administration. United states government's open data. <https://data.gov/>. Accessed: 2025-02-11.
- [2] Jetstream2 Documentation. Configuring a vm to use manila shares. <https://docs.jetstream-cloud.org/general/manilaVM/>. Accessed: 2025-02-04.
- [3] David Y Hancock, Jeremy Fischer, John Michael Lowe, Winona Snapp-Childs, Marlon Pierce, Suresh Marru, J Eric Coulter, Matthew Vaughn, Brian Beck,

- Nirav Merchant, et al. Jetstream2: Accelerating cloud computing via jetstream. In *Practice and Experience in Advanced Research Computing*, pages 1–8. 2021.
- [4] Thomas Kluyver, Benjamin Ragan-Kelley, Fernando Pérez, Brian Granger, Matthias Bussonnier, Jonathan Frederic, Kyle Kelley, Jessica Hamrick, Jason Grout, Sylvain Corlay, et al. Jupyter notebooks—a publishing format for reproducible computational workflows. In *Positioning and power in academic publishing: Players, agents and agendas*, pages 87–90. IOS press, 2016.
- [5] Matthew J Menne, Imke Durre, Russell S Vose, Byron E Gleason, and Tamara G Houston. An overview of the global historical climatology network-daily database. *Journal of atmospheric and oceanic technology*, 29(7):897–910, 2012.
- [6] Deborah Nolan and Duncan Temple Lang. Computing in the statistics curricula. *The American Statistician*, 64(2):97–107, 2010.
- [7] Seref Sagiroglu and Duygu Sinanc. Big data: A review. In *2013 international conference on collaboration technologies and systems (CTS)*, pages 42–47. IEEE, 2013.
- [8] Internal Revenue Service. Form 990 series downloads. <https://www.irs.gov/charities-non-profits/form-990-series-downloads>. Accessed: 2025-01-28.
- [9] Amazon Web Services. Amazon open data. <https://registry.opendata.aws/>. Accessed: 2025-02-11.

- [10] Amazon Web Services. Instance type explorer. <https://aws.amazon.com/ec2/instance-explorer>. Accessed: 2025-02-11.
- [11] Indiana University. Jetstream2 main website. <https://jetstream-cloud.org/index.html>. Accessed: 2025-02-04.
- [12] Hadley Wickham, Mine Çetinkaya-Rundel, and Garrett Grolemund. *R for data science*. " O'Reilly Media, Inc.", 2023.
- [13] Andy B Yoo, Morris A Jette, and Mark Grondona. Slurm: Simple linux utility for resource management. In *Workshop on job scheduling strategies for parallel processing*, pages 44–60. Springer, 2003.